



# Serverless

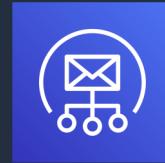
AWS Certified Solutions Architect – Associate

Harun Ur Rashid – Cloud Infrastructure Architect  
27<sup>th</sup> of August 2021



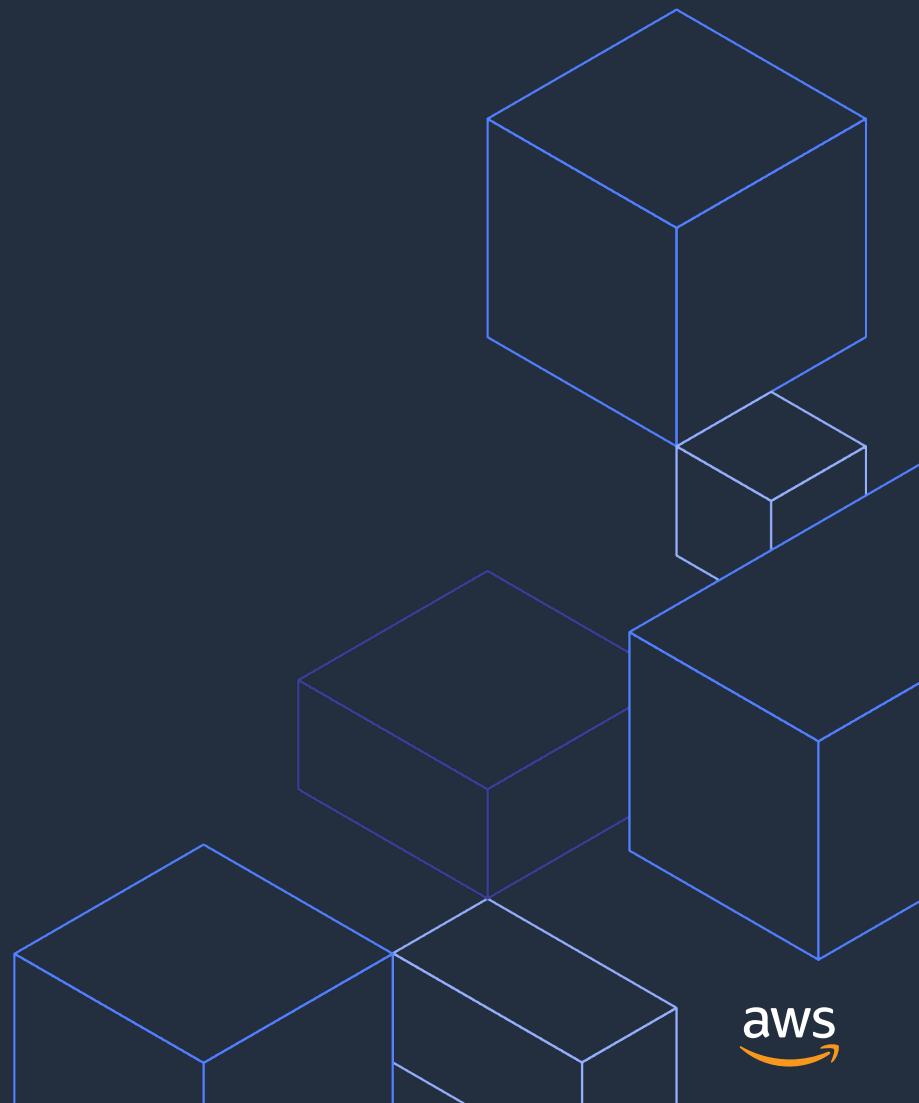
# Agenda

- What is *Serverless*?
- AWS Lambda and AWS Step Functions
- Amazon Simple Queue Service (SQS)
- Amazon Simple Notification Service (SNS)
- Amazon API Gateway
- Amazon Simple Email Service (SES)
- Amazon Elastic Container Service (ECS)
- AWS Well Architected Framework

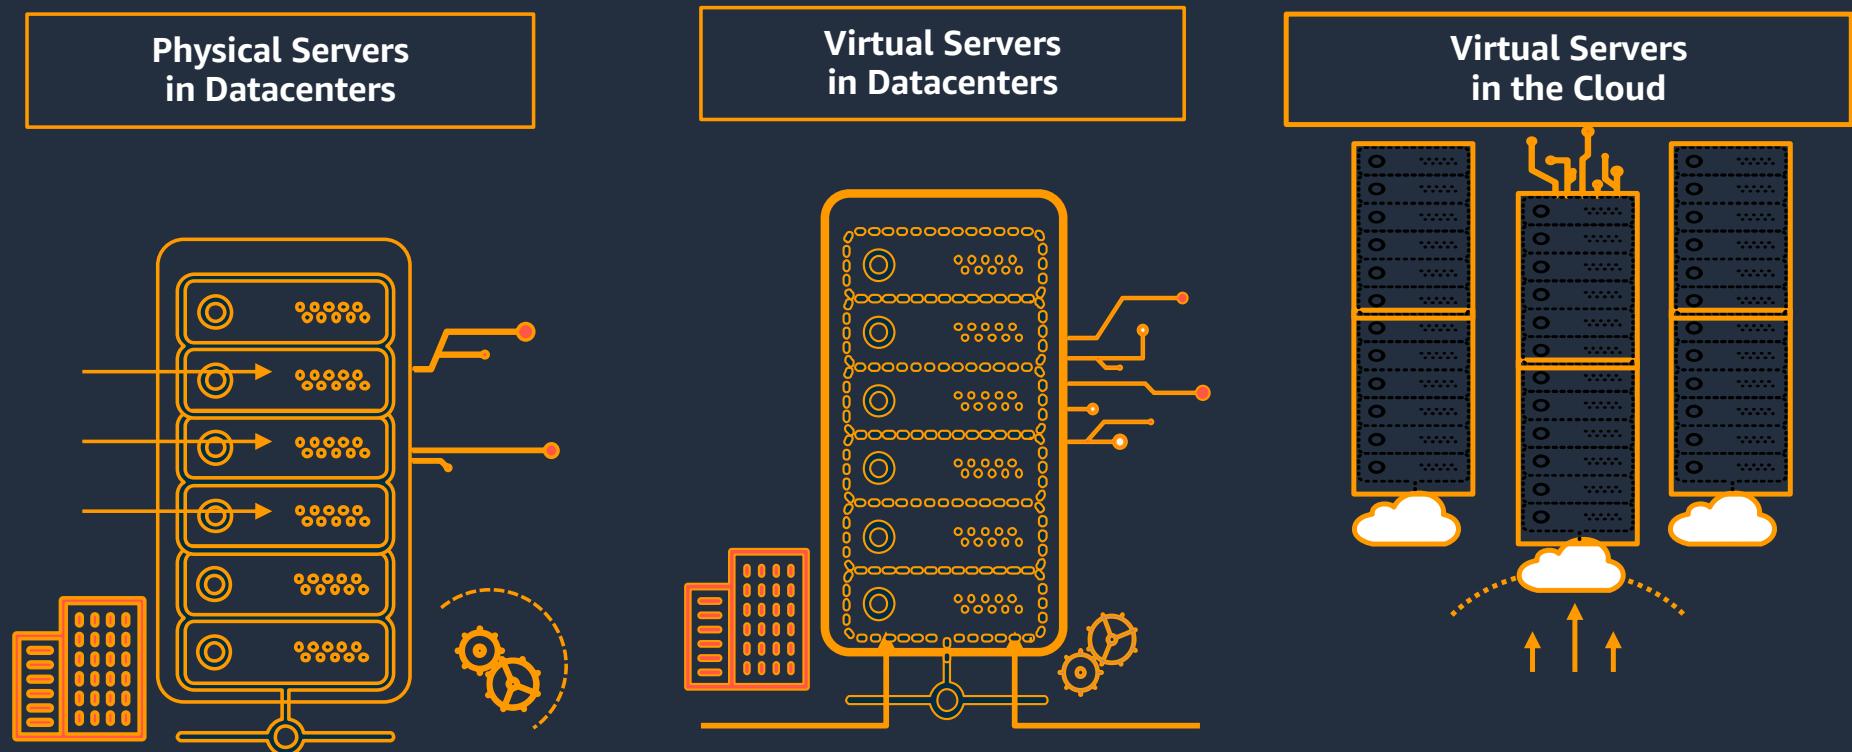


# What is Serverless?

© 2021, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



# Evolution of computing



# What is *Serverless*?



No infrastructure provisioning,  
no management



Automatic scaling

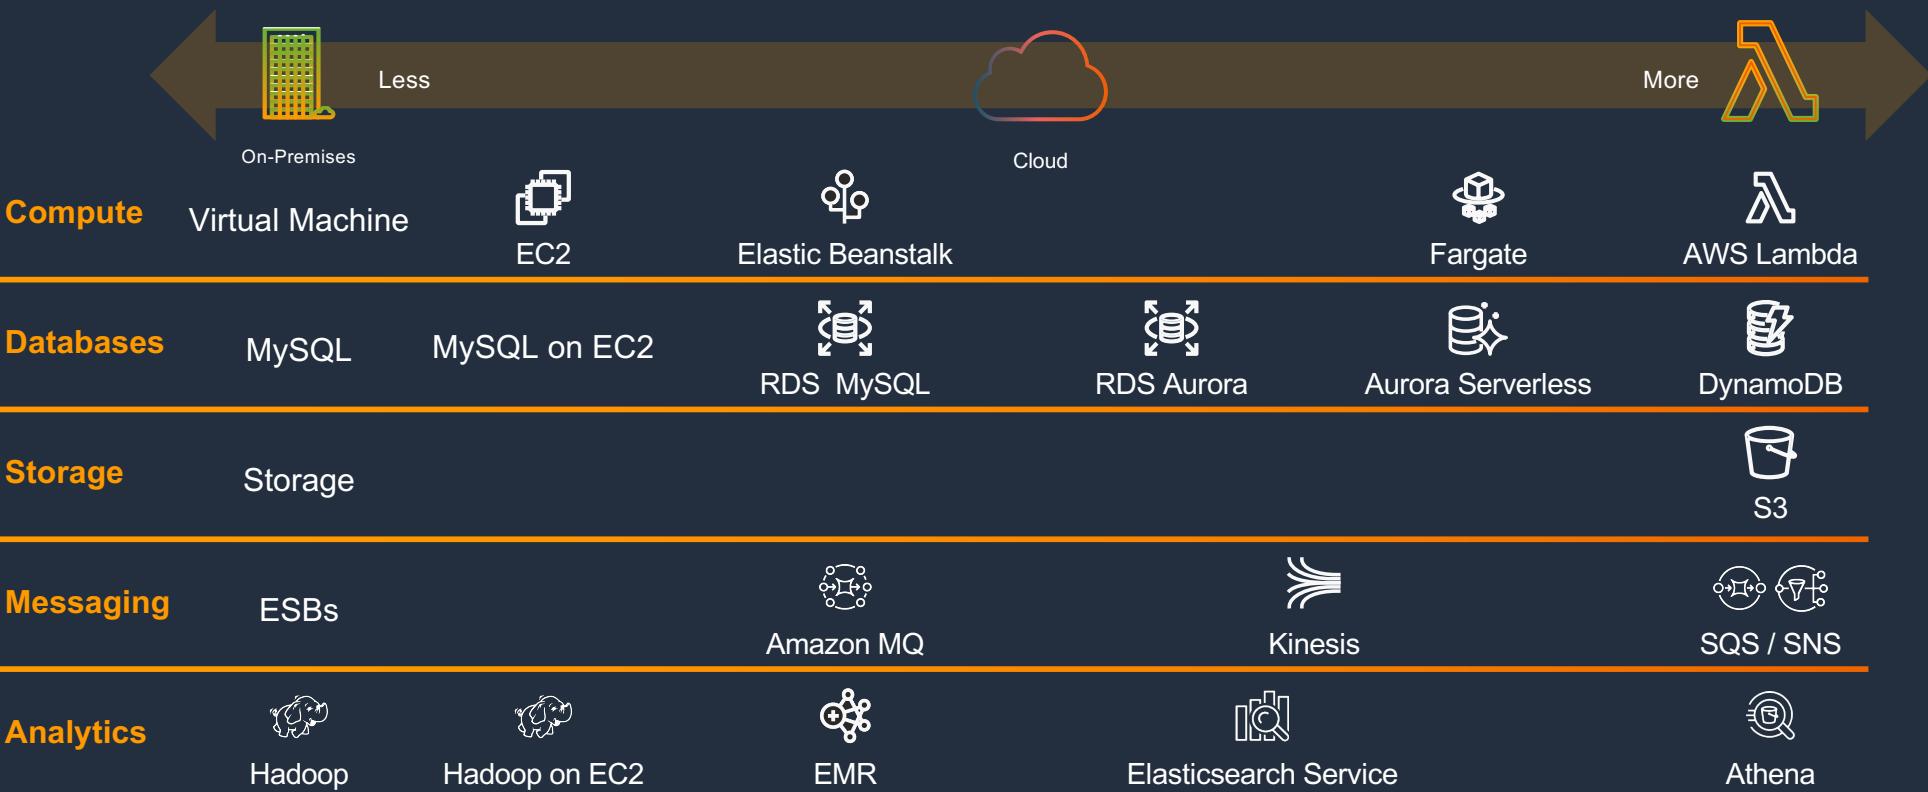
Pay for value



Highly available and secure



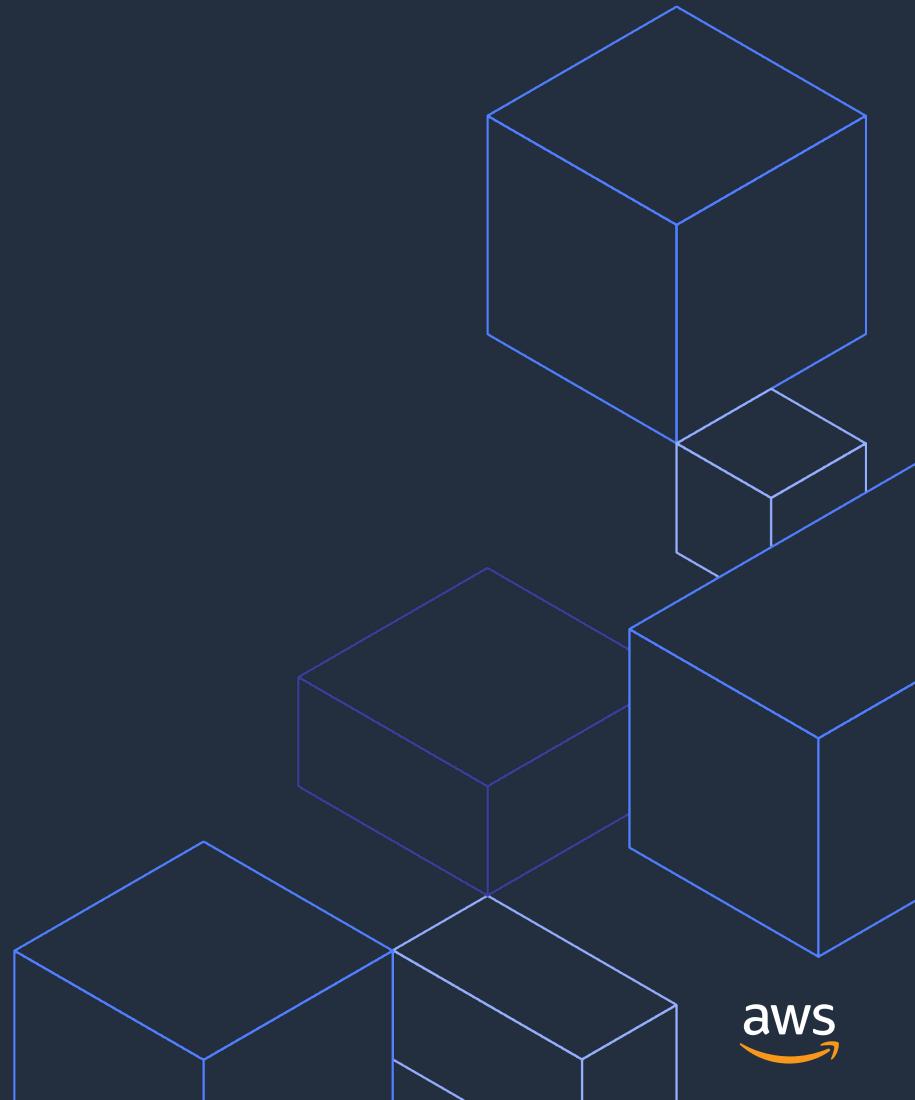
# AWS operational responsibility models



# AWS Lambda



© 2021, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



# Introduction to AWS Lambda

- “Function-as-a-Service”
- Run code without provisioning or managing servers
  - Java, Go, PowerShell, Node.js, C#, Python, and Ruby
  - Runtime API which supports additional programming languages
- Pay only for the compute time you consume
- Automatically runs your code with high availability
- Scale with usage



# How Lambda works

## Invoked in response to events

- Changes in data
- Changes in state



S3 event notifications



DynamoDB Streams



Kinesis events



SNS events



CloudTrail events



Cognito events



Custom events



Custom events

Author in familiar language  
using any libraries; Execute  
only when needed,  
automatic scale



“Lambda  
functions”

Access any service,  
including your own

Any custom



Any AWS



Such as...

SNS

DynamoDB Lambda



Redshift



Kinesis

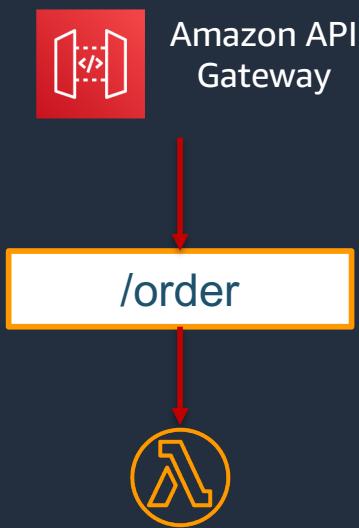


S3

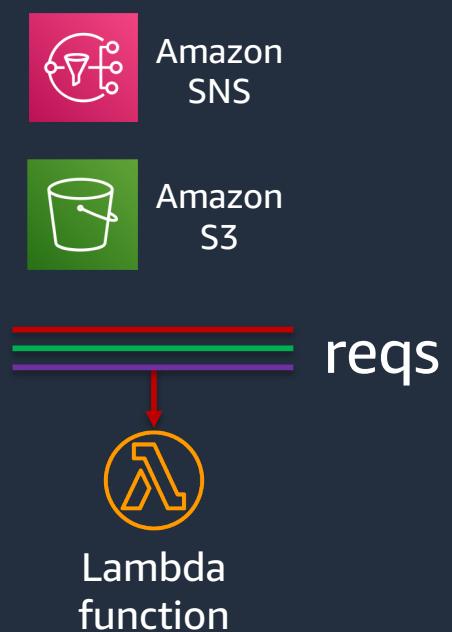


# Lambda execution models

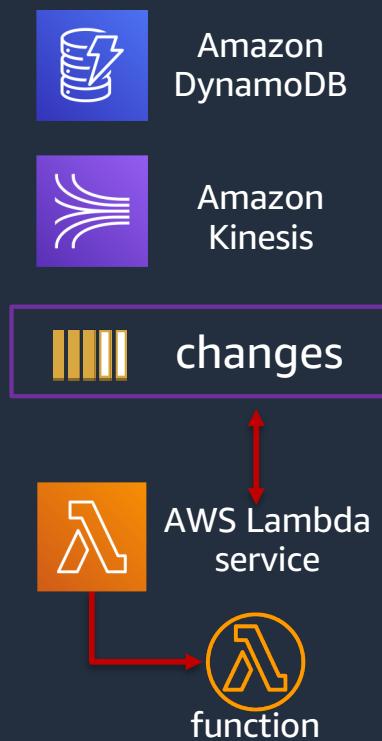
## Synchronous (push)



## Asynchronous (event)



## Stream (Poll-based)



# AWS event sources that can trigger Lambda

## Synchronous (push):

Elastic Load Balancing (Application Load Balancer)  
Amazon Cognito  
Amazon Lex  
Amazon Alexa  
Amazon API Gateway  
Amazon CloudFront (Lambda@Edge)  
Amazon Kinesis Data Firehose  
Amazon Simple Storage Service Batch

## Stream (poll):

Amazon DynamoDB  
Amazon Kinesis  
Amazon Managed Streaming for Apache Kafka  
Amazon Simple Queue Service

## Asynchronous (event):

Amazon Simple Storage Service  
Amazon Simple Notification Service  
Amazon Simple Email Service  
AWS CloudFormation  
Amazon CloudWatch Logs  
Amazon CloudWatch Events  
AWS CodeCommit  
AWS Config  
AWS IoT  
AWS IoT Events  
AWS CodePipeline

# Anatomy of a Lambda function

## Handler function

Function to be executed upon invocation

## Event object

Data sent during Lambda function invocation

## Context object

Methods available to interact with runtime information (request ID, log group, more)

```
import json

def lambda_handler(event, context):
    # TODO implement
    return {
        'statusCode': 200,
        'body': json.dumps('Hello World!')
    }
```

# Anatomy of a Lambda function – Example

## Input

```
{  
  "my_number": 5  
}
```



## Response

```
Response:  
{  
  "statusCode": 200,  
  "body": "25"  
}
```

```
import json  
  
def lambda_handler(event, context):  
    my_number = event["my_number"]  
    my_number_squared = my_number * my_number  
    return {  
        'statusCode': 200,  
        'body': json.dumps(my_number_squared)  
    }
```

# Lambda Layers



Lets functions easily share code: Upload layer once, reference within any function

Promote separation of responsibilities, lets developers iterate faster on writing business logic

Built in support for secure sharing by ecosystem

# Additional Lambda configurations

- Assignable memory in 64MB increments, from 128MB to 3GB
- Configured to run up to 15 minutes per execution
- Environment variables
- Reserved and provisioned concurrency
- Amazon VPC, Security Groups, and RDS Proxy
- Stateless /tmp storage, and stateful filesystem using Amazon EFS
- Publish new versions of functions, and alias them
- AWS Serverless Application Model (AWS SAM)
- Lambda @ Edge

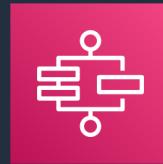
## Poll#1

- Which of the following services can invoke a Lambda function synchronously (with functionality built-in with the invoking service)?
- A: IAM
- B: Amazon Lex
- C: Amazon Kinesis Data Firehose
- D: API Gateway
- E: S3

## Poll#2

- What is the maximum amount of memory you can allocate to an AWS Lambda function?
- A: 128MB
- B: 512MB
- C: 1GB
- D: 3GB

# AWS Step Functions



© 2021, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



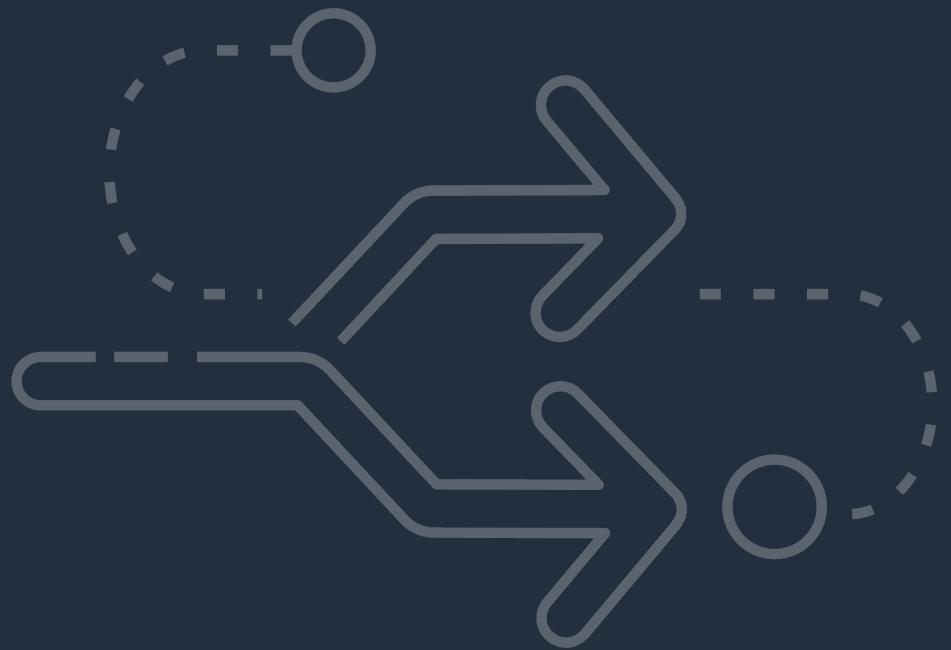
# State machine

Describes a **collection of computational steps** split into discrete states

Has **one starting state and always one active state** (while executing)

The **active state** receives input, takes some action, and generates output

**Transitions between states** are based on state outputs and rules that we define



# AWS Step Functions

Fully-managed state machines on AWS

Resilient workflow automation

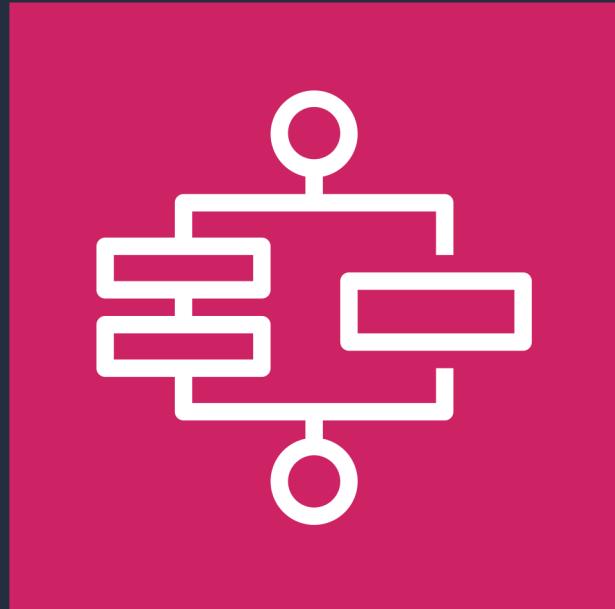
Built-in error handling

Powerful AWS service integration

First-class support for integrating with  
your own services

Auditable execution history and visual monitoring

© 2021, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



# How AWS Step Functions work



The **workflows** you build with Step Functions are called **state machines**, and **each step** of your workflow is called a **state**.



When you execute your state machine, **each move** from one state to the next is called a **state transition**.



You can **reuse components**, easily edit the sequence of steps or swap out the code called by task states as your needs change.

# Amazon State Language

<https://states-language.net/spec.html>

```
{  
  "Comment": "A simple minimal example",  
  "StartAt": "Hello World",  
  "States": {  
    "Hello World": {  
      "Type": "Task",  
      "Resource": "arn:aws:lambda...HelloWorld",  
      "End": true  
    },  
    [...]  
  }  
}
```



# Example workflow: opening an account



Tasks



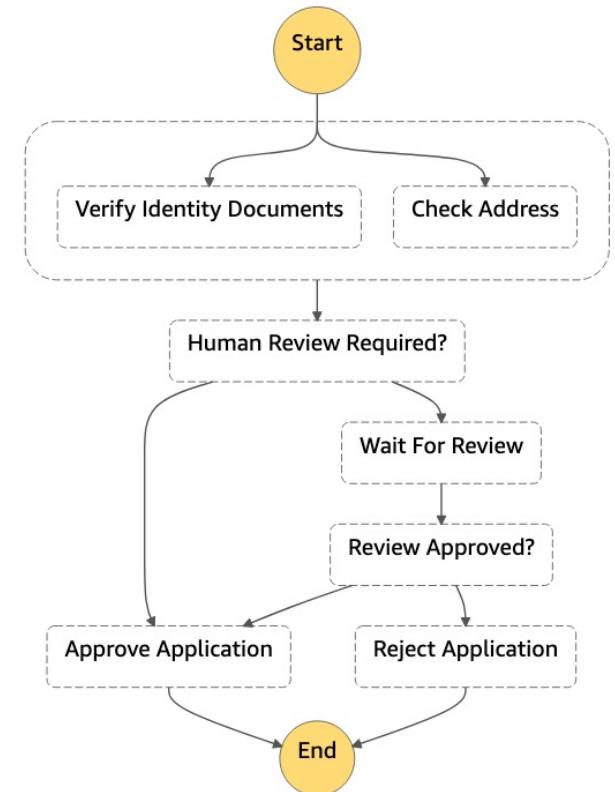
Parallel Steps



Branching Choice



Wait for a callback



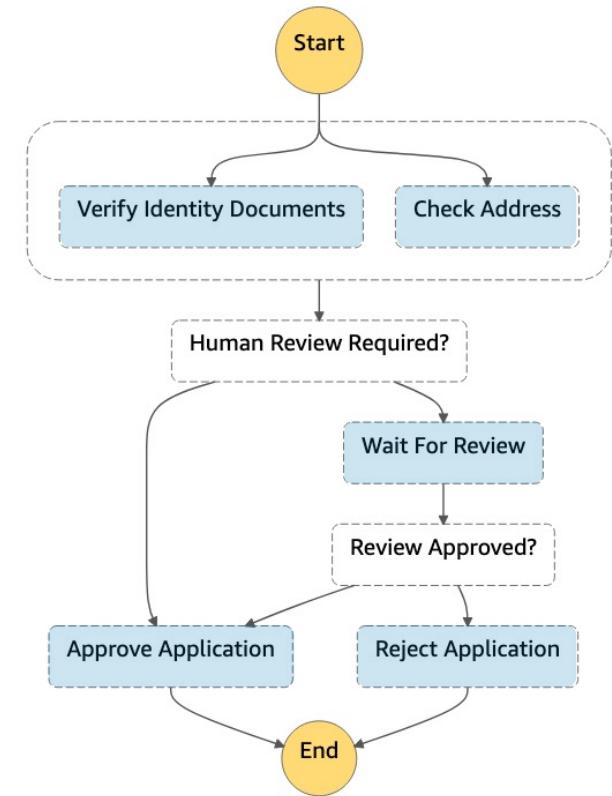


## Performing a task

Call an AWS Lambda Function

Wait for a polling worker to perform an activity

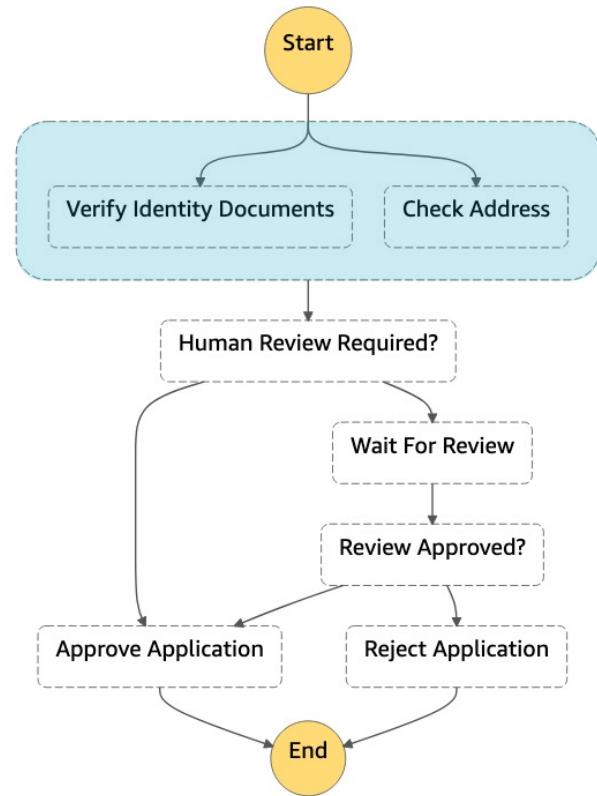
Pass parameters to an API of an integrated AWS Service



## Executing branches in parallel

Contains an array of state machines branches to execute in parallel

Outputs an array of outputs from each state machine in its branches



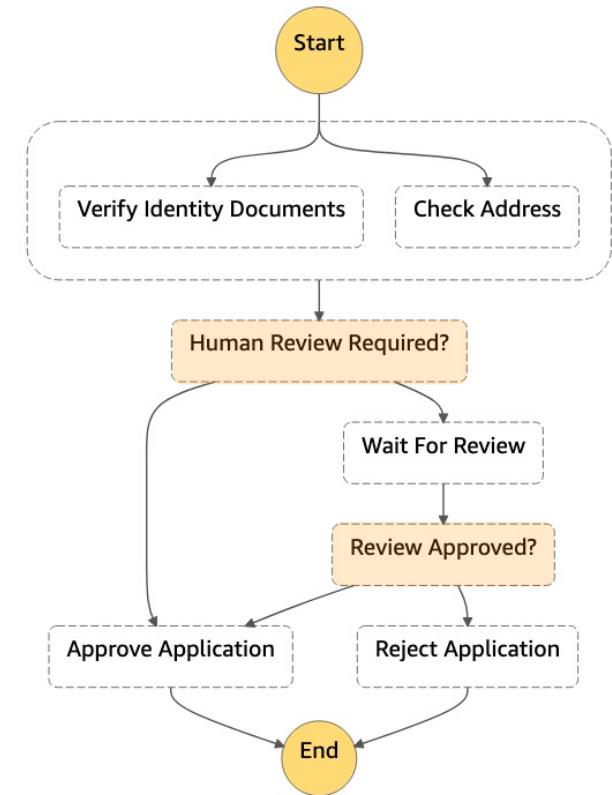


## Making a choice

Like a switch statement in programming

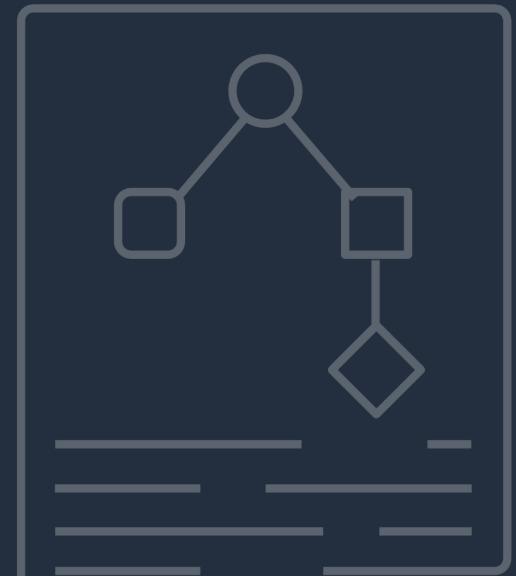
Inspects an array of *choice* expressions, comparing variables to values

Determines which state to transition to next



# State types

<b>Task</b>	Execute work
<b>Choice</b>	Add branching logic
<b>Wait</b>	Add a timed delay
<b>Parallel</b>	Execute branches in parallel
<b>Map</b>	Process each of an input array's items with a state machine
<b>Succeed</b>	Signal a successful execution and stop
<b>Fail</b>	Signal a failed execution and stop
<b>Pass</b>	Pass input to output



# AWS Step Functions service integrations



AWS  
Lambda



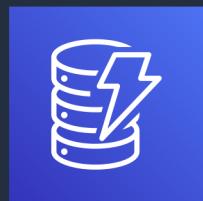
Amazon  
Elastic Container Service



AWS  
Batch



AWS CodeBuild



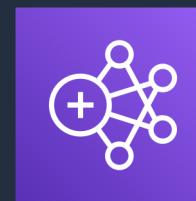
Amazon  
DynamoDB



AWS  
Glue



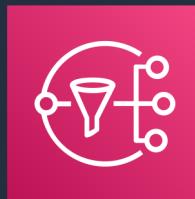
Amazon  
SageMaker



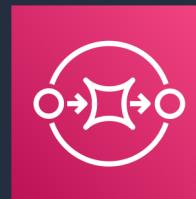
Amazon EMR



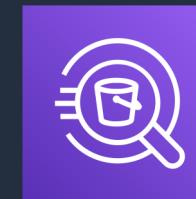
AWS  
Step Functions



Amazon  
Simple Notification Service



Amazon  
Simple Queue Service



Amazon Athena



## Poll#3

- AWS Step Functions can orchestrate multiple AWS Glue Jobs – True or False?
- A: True
- B: False

# Amazon Simple Queue Service (SQS)



© 2021, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



# What is a “queue”?



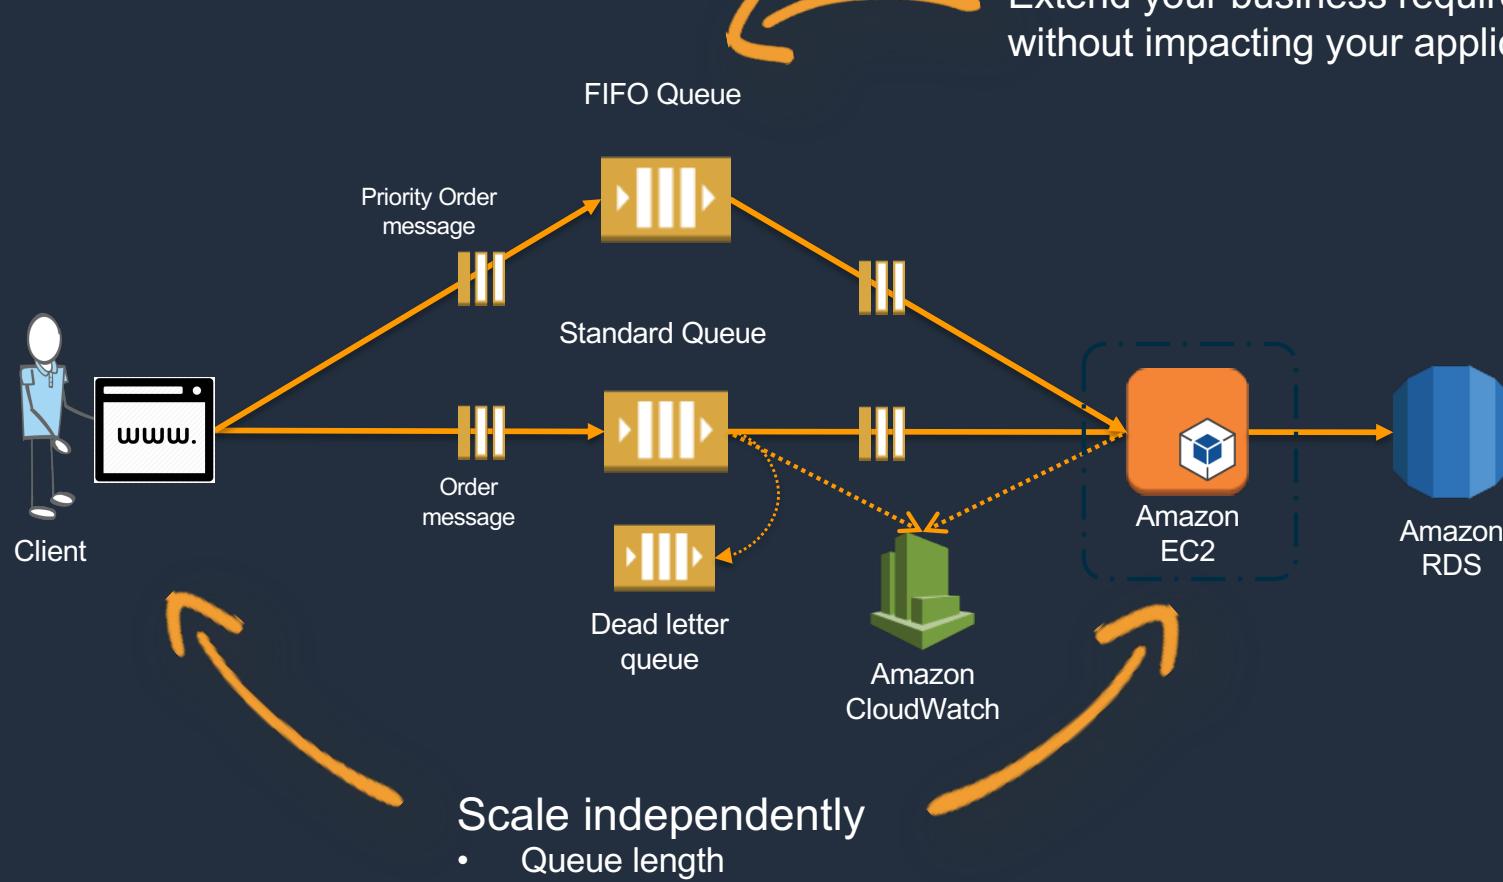
# Amazon SQS overview



- Fast, reliable, scalable, fully managed **queue service**
- **Send, store, and receive** messages between components
- Consumer **polls** for messages
- AWS **SDK** or Java Message Service (**JMS**) APIs
- Pay for **what you use**

# Amazon SQS

Extend your business requirements  
without impacting your application logic



© 2021, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



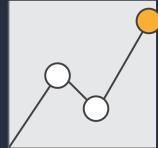
# Key features



- Persistent message queue with high **durability** and **availability**
- Messages are stored across **multiple AZs**
- Messages retained until deleted— **up to 14 days**
- Nearly **unlimited throughput**



# Key features



- Amazon CloudWatch metrics and alerts

`1001  
0010  
0101  
1011`

- Message payloads up to 256 KB (without Extended Client Library for Java)



- Message batching to increase throughput and reduce cost



- Secure: uses AWS Identity & Access Management (IAM) and HTTPS/TLS

# SQS queue types

## Standard Queues

- At least once delivery
- Best-effort ordering
- Nearly unlimited transaction rate

## FIFO Queues

- First-in First-out delivery preserving message ordering
- Exactly once processing
- 300 transactions per second
- Messages groups for multiple ordered streams

## Poll#4

- Which of the following SQS Queue preserve String Ordering in message processing?
- A: Standard Queues
- B: FIFO Queues

# Amazon Simple Notification Service (SNS)



© 2021, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



# What is “pub-sub”?



# What is Amazon SNS?

- Amazon **SNS** is a web service that coordinates and manages the delivery or sending of messages to subscribing endpoints or clients
- Two types of clients—**publishers** and **subscribers**—also referred to as **producers** and **consumers**
- Publishers communicate **asynchronously** with subscribers by producing and sending a message to a topic
- Subscribers like web servers, email addresses, Amazon SQS queues, AWS Lambda functions consume or receive the message or notification through:
  - Amazon **SQS**, **HTTP/S**, **email**, **SMS**, **Lambda**



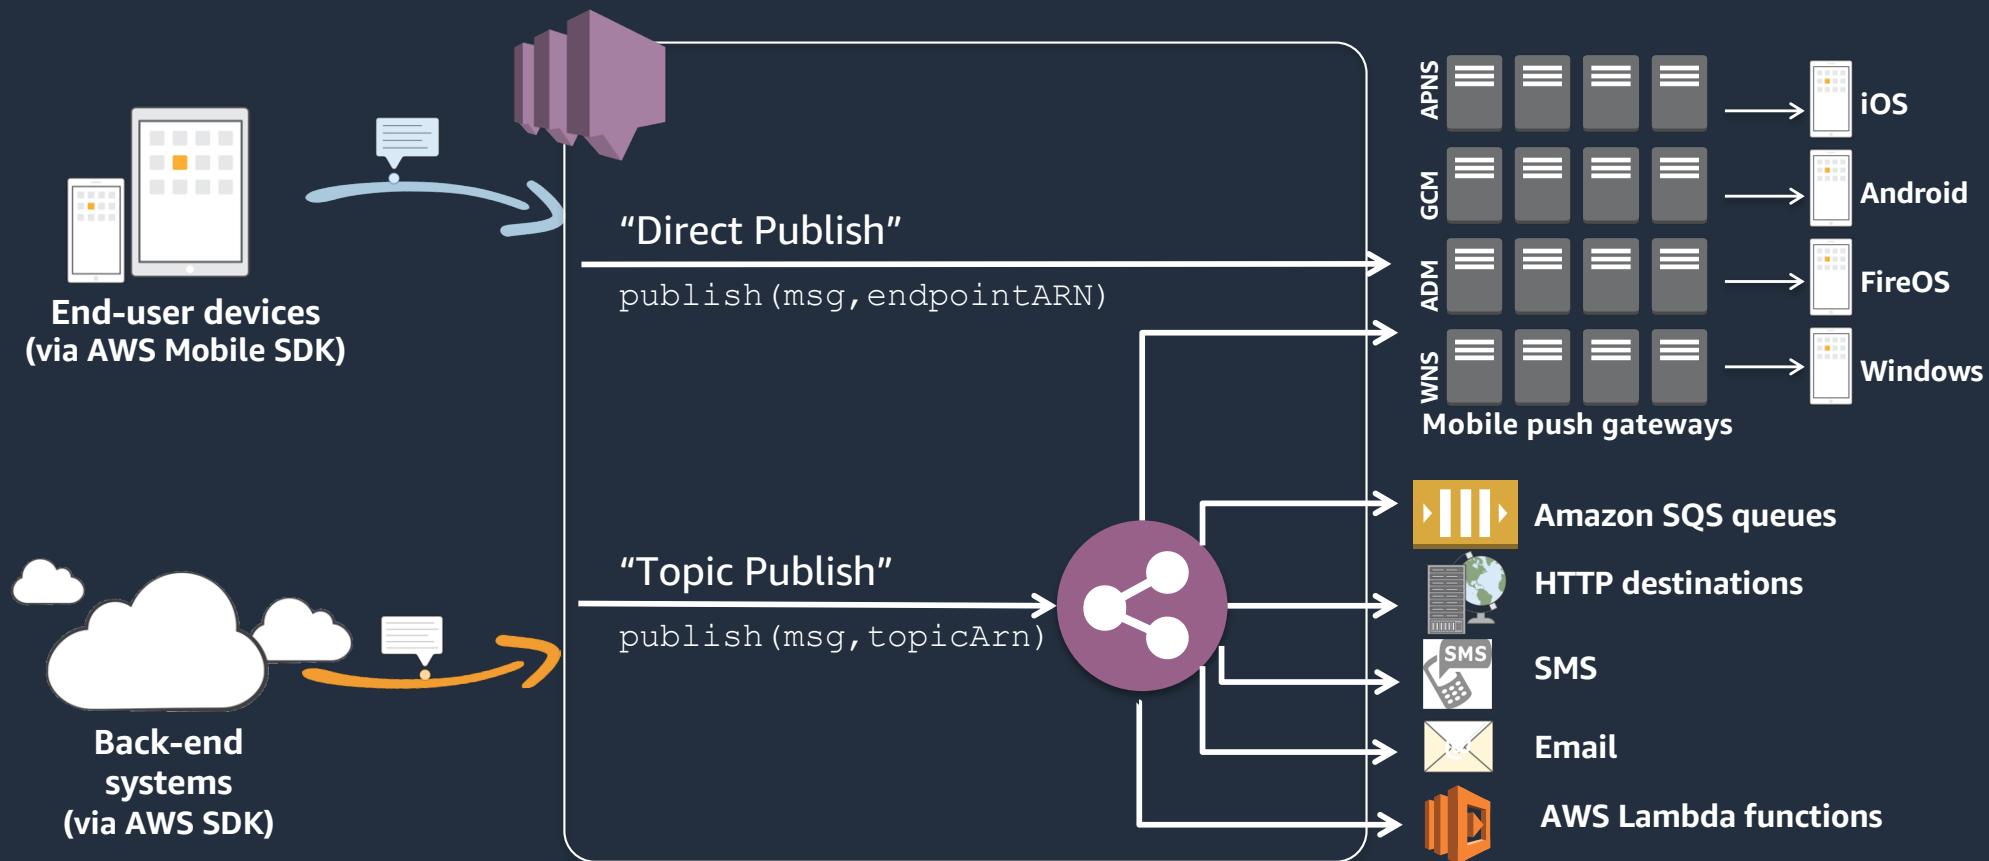
# How Amazon SNS works

- 1) Create a Topic
- 2) Subscribe to a Topic
- 3) Publish to a Topic

See full API list:  
<http://docs.aws.amazon.com/sns/latest/api/Welcome.html>



# One API - Multiple destination types



# SQS vs SNS



SQS



SNS

- Queue
- Pull
- Decoupling two applications
- Messages are persisted in the queue for a fixed duration

- Topic (Pub/Sub)
- Push
- Fanout from one to many applications
- No persistence (but can retry)

## Poll#5

- Which delivery mechanism is NOT supported by Amazon Simple Notification Service (SNS)?
- A: Email
- B: HTTP(S)
- C: Short Message Service (SMS)
- D: Really Simple Syndication (RSS)

# Amazon API Gateway

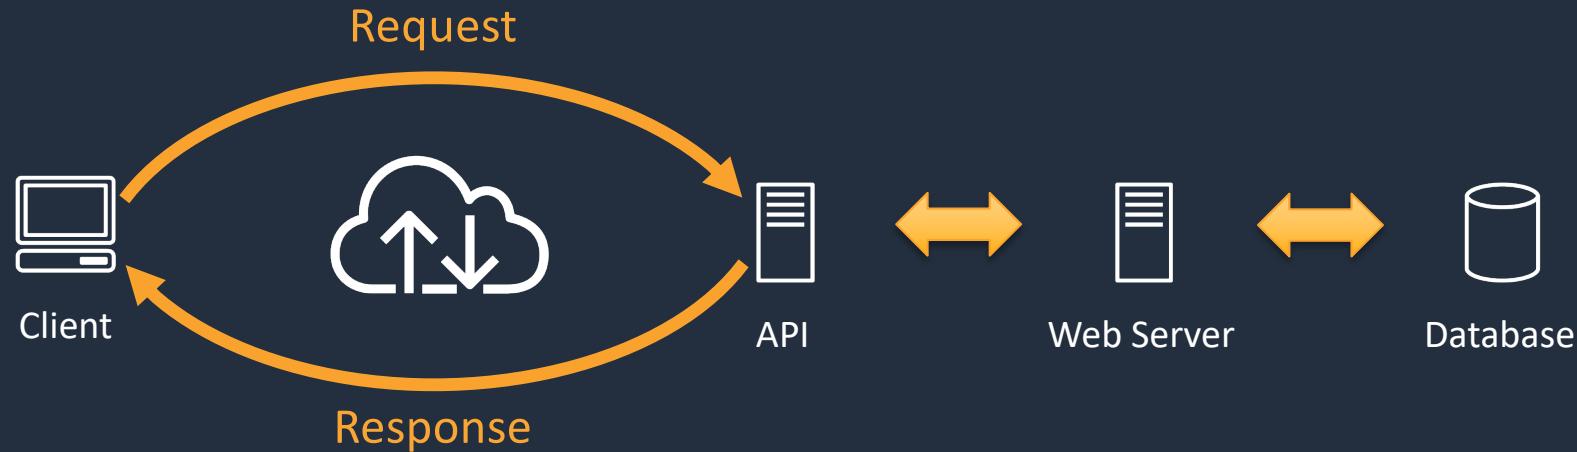


© 2021, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



# Application Programming Interface (API)

“ In building applications, an API simplifies programming by abstracting the underlying implementation and only exposing objects or actions the developer needs.



Web-based companies and services offer APIs for developers to use, such as:

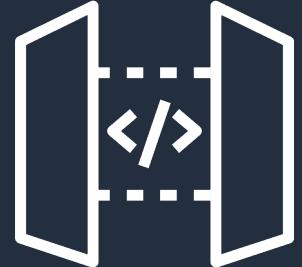
- Social Networks – Facebook, Twitter, etc
- Payment Processing – Amazon Pay, PayPal, etc

# Amazon API Gateway

API Gateway is a fully managed service that makes it easy for developers to create, publish, maintain, monitor, and secure APIs at any scale. It frees you from the operational burden of implementation, offers reliable network protection, and centralizes authorization decisions within policies so bugs and code concerns are minimized.

*It also enables you to:*

- Host multiple versions and stages of your APIs
- Create and distribute API Keys to developers
- Throttle and monitor requests to protect your backend
- Leverage signature version 4 to authorize access to APIs
- Perform Request / Response data transformation and API mocking
- Reduce latency and DDoS protection through CloudFront
- Store API responses through managed caches
- Generate SDKs for Java, JavaScript, Java for Android, Objective-C or Swift for iOS, and Ruby



# Supported protocols

RESTful: HTTP APIs & REST APIs



- Request / Response
- HTTP Methods like GET, POST, etc
- Short-lived communication
- Stateless

WebSocket APIs



- Serverless WebSocket
- 2 way communication channel
- Long-lived communication
- Stateful

# Types of APIs

## RESTful APIs

[HTTP APIs](#) are the cheapest, fastest, best choice for building APIs that only require API proxy functionality. For APIs that require API proxy functionality and management features in a single solution, API Gateway also offers [REST APIs](#).

### Regional (*Available with all types*)

- Recommended API type for general use cases
- Designed for building APIs for clients in the same region

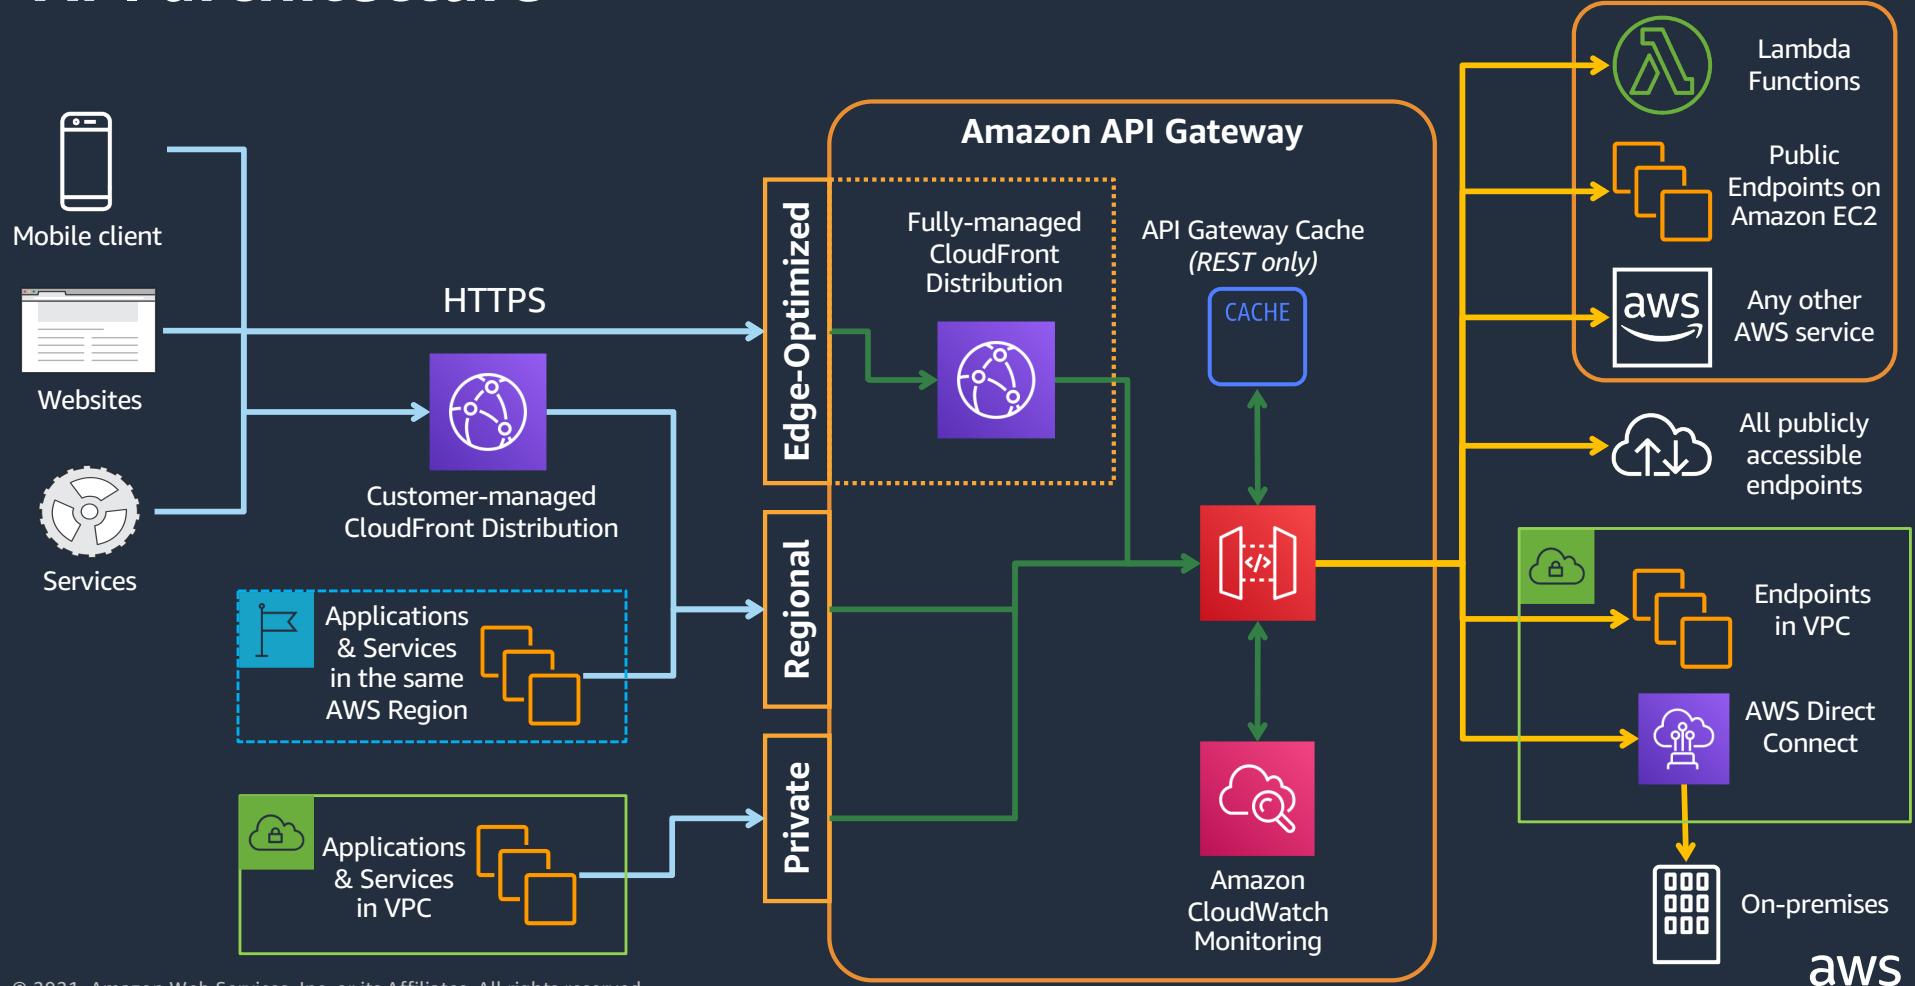
### Edge-Optimized (*Available with REST APIs*)

- Uses CloudFront to reduce TLS connection overhead (reduces roundtrip time)
- Designed for a globally distributed clients

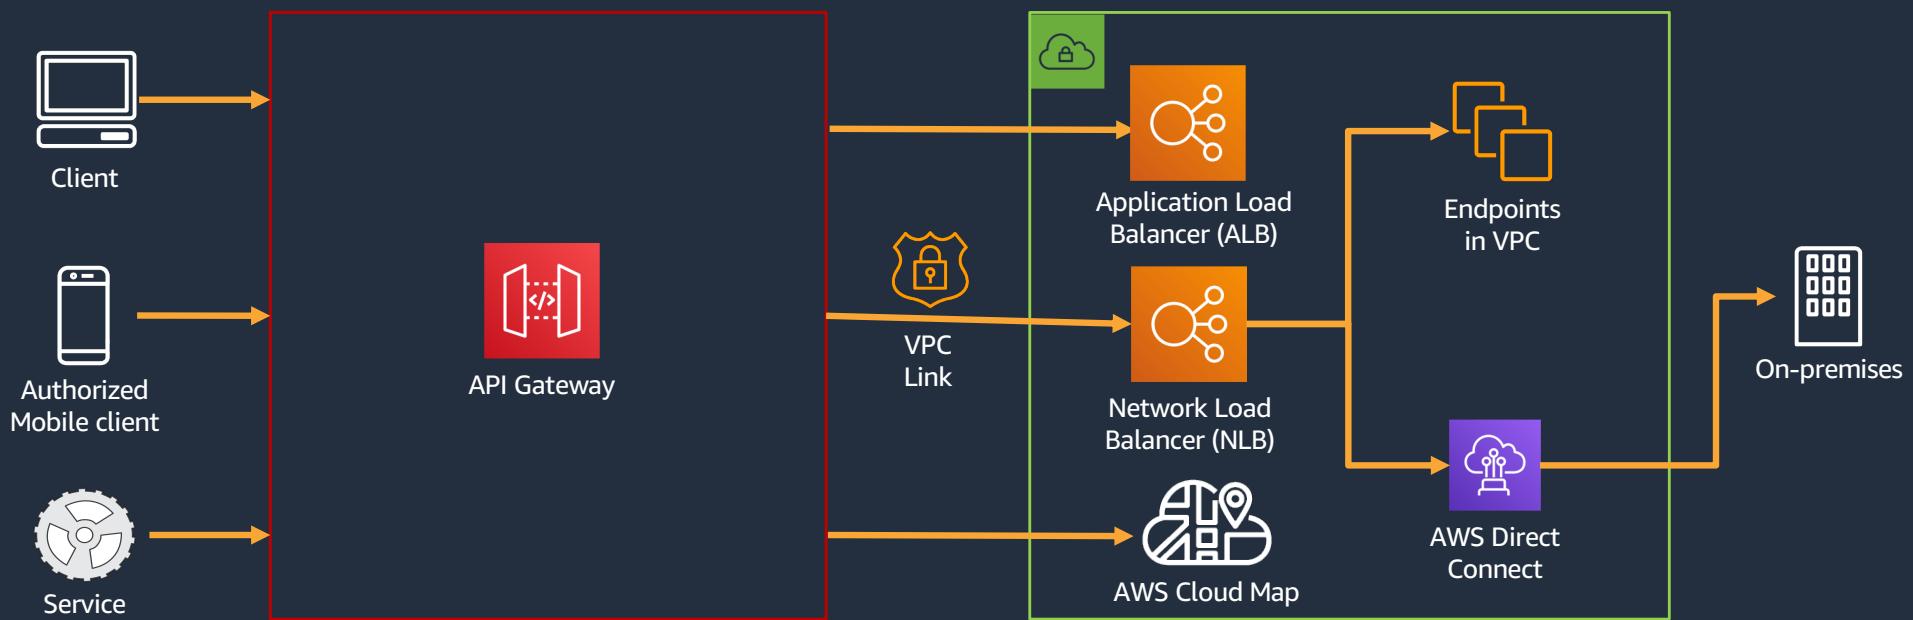
### Private (*Available with REST APIs*)

- Only accessible from within VPC (and networks connected to VPC)
- Designed for building APIs used internally or by private microservices

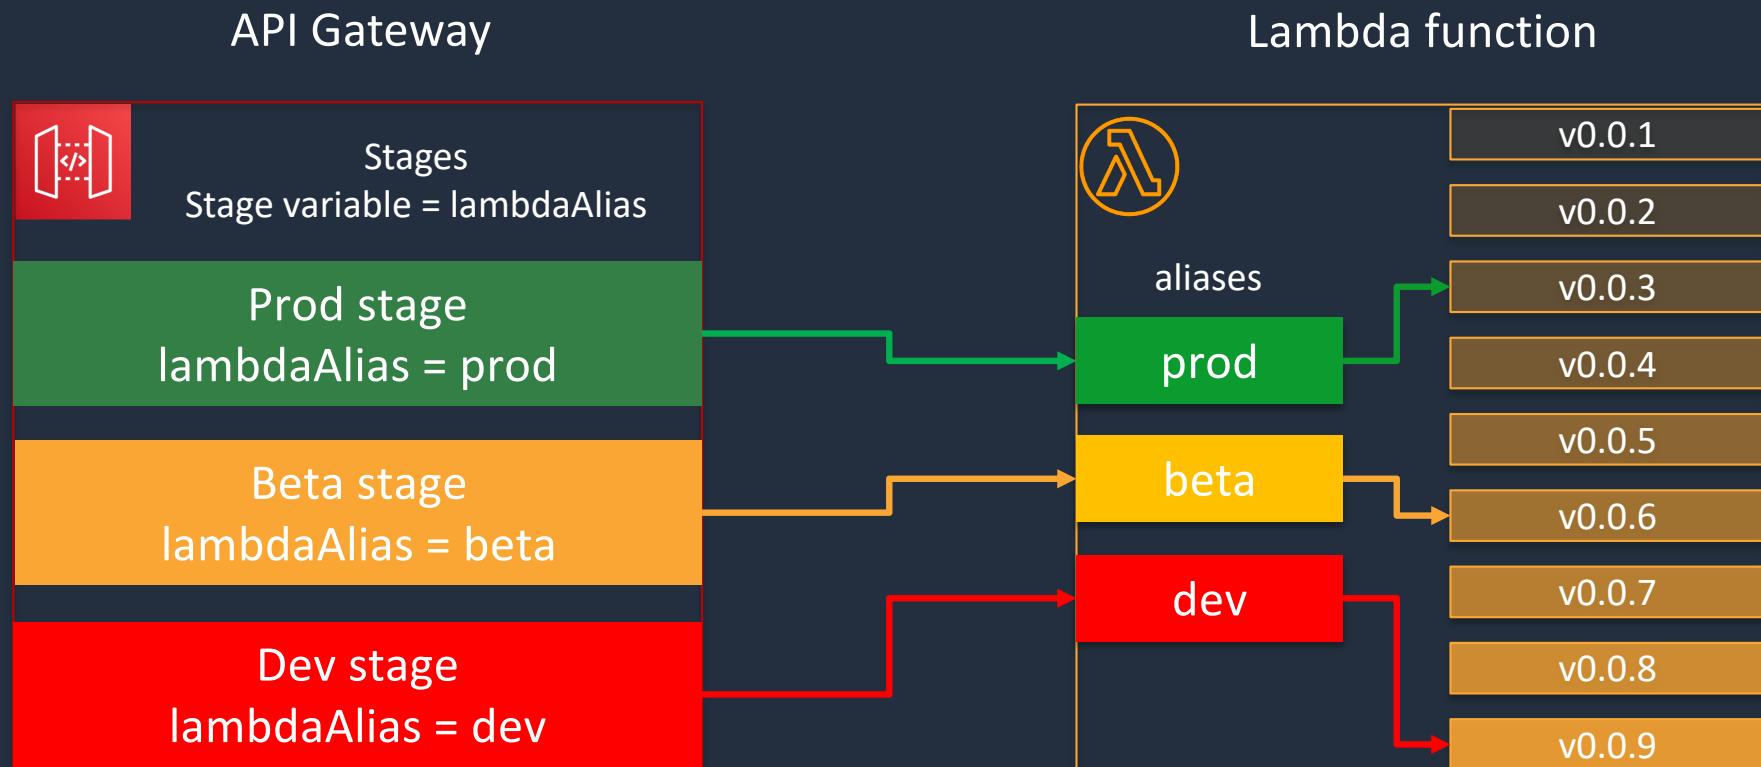
# API architecture



# VPC Links (Private Integrations)



# Staging



# Custom domains

~~https://12345.execute-api.us-east-1.amazonaws.com/prod~~

**https://mydomain.com/api-one**

- Supports HTTP, REST, and WebSocket APIs
- SSL Certs managed through ACM
- Supports multiple domains through base path mapping

## Poll#6

- Which of the following Component/Feature of API Gateway enabled Private Integration with AWS Backend Resources?
- A: Stage
- B: Custom Domain
- C: Method
- D: Resource
- E: VPC Links
- F: APIKEY

# Amazon Simple Email Service



© 2021, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



Amazon SES is a flexible, affordable, and highly-scalable email sending and receiving service for businesses and developers



High performance



High deliverability



Scalable



Reliable



Configurable

# Amazon Elastic Container Service (ECS)



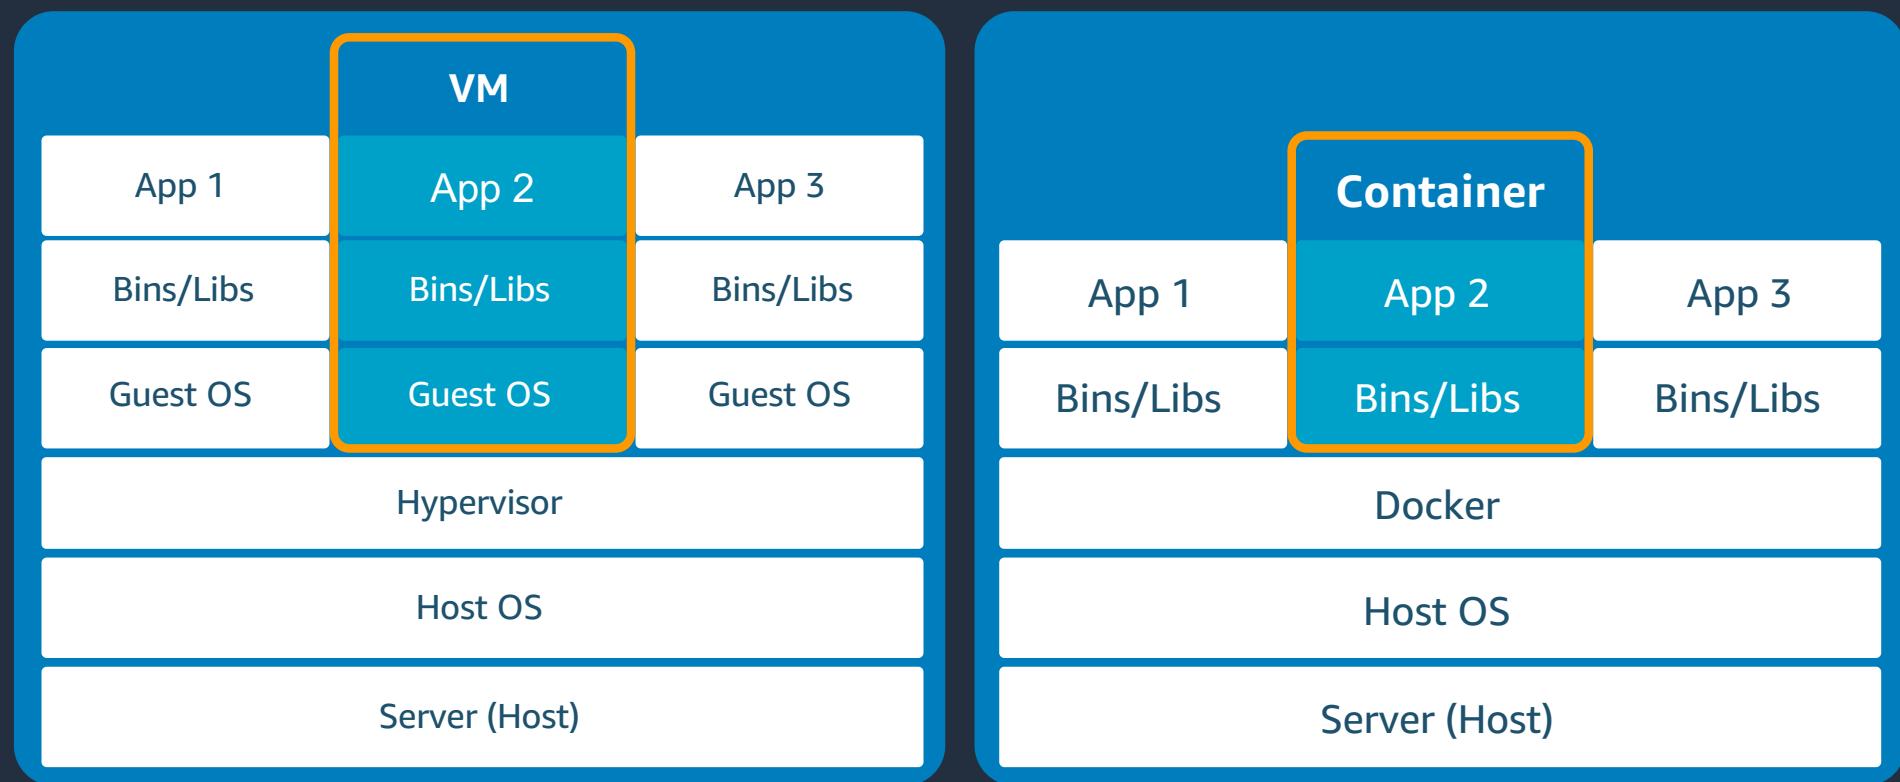
© 2021, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



# What are containers?



# What are containers?



# AWS container services landscape

## Management

Deployment, Scheduling,  
Scaling & Management of  
containerized applications

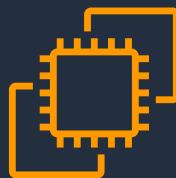


Amazon Elastic Container Service

Amazon Elastic Kubernetes service

## Hosting

Where the containers run



Amazon EC2



AWS Fargate

## Image Registry

Container Image Repository



Amazon Elastic Container Registry

## Poll#7

- Which of the following AWS Container as a Service Platform supports Serverless Architecture/Deployment
- A: Amazon Elastic Container Service
- B: AWS Fargate
- C: Amazon Elastic Kubernetes Service

# Well-Architected Framework



© 2021, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



# Pillars



# General principles

Stop guessing your capacity needs

---

Test systems at production scale

---

Automate to make architectural experimentation easier

---

Allow for evolutionary architectures

---

Drive architectures using data

---

Improve through game days

---

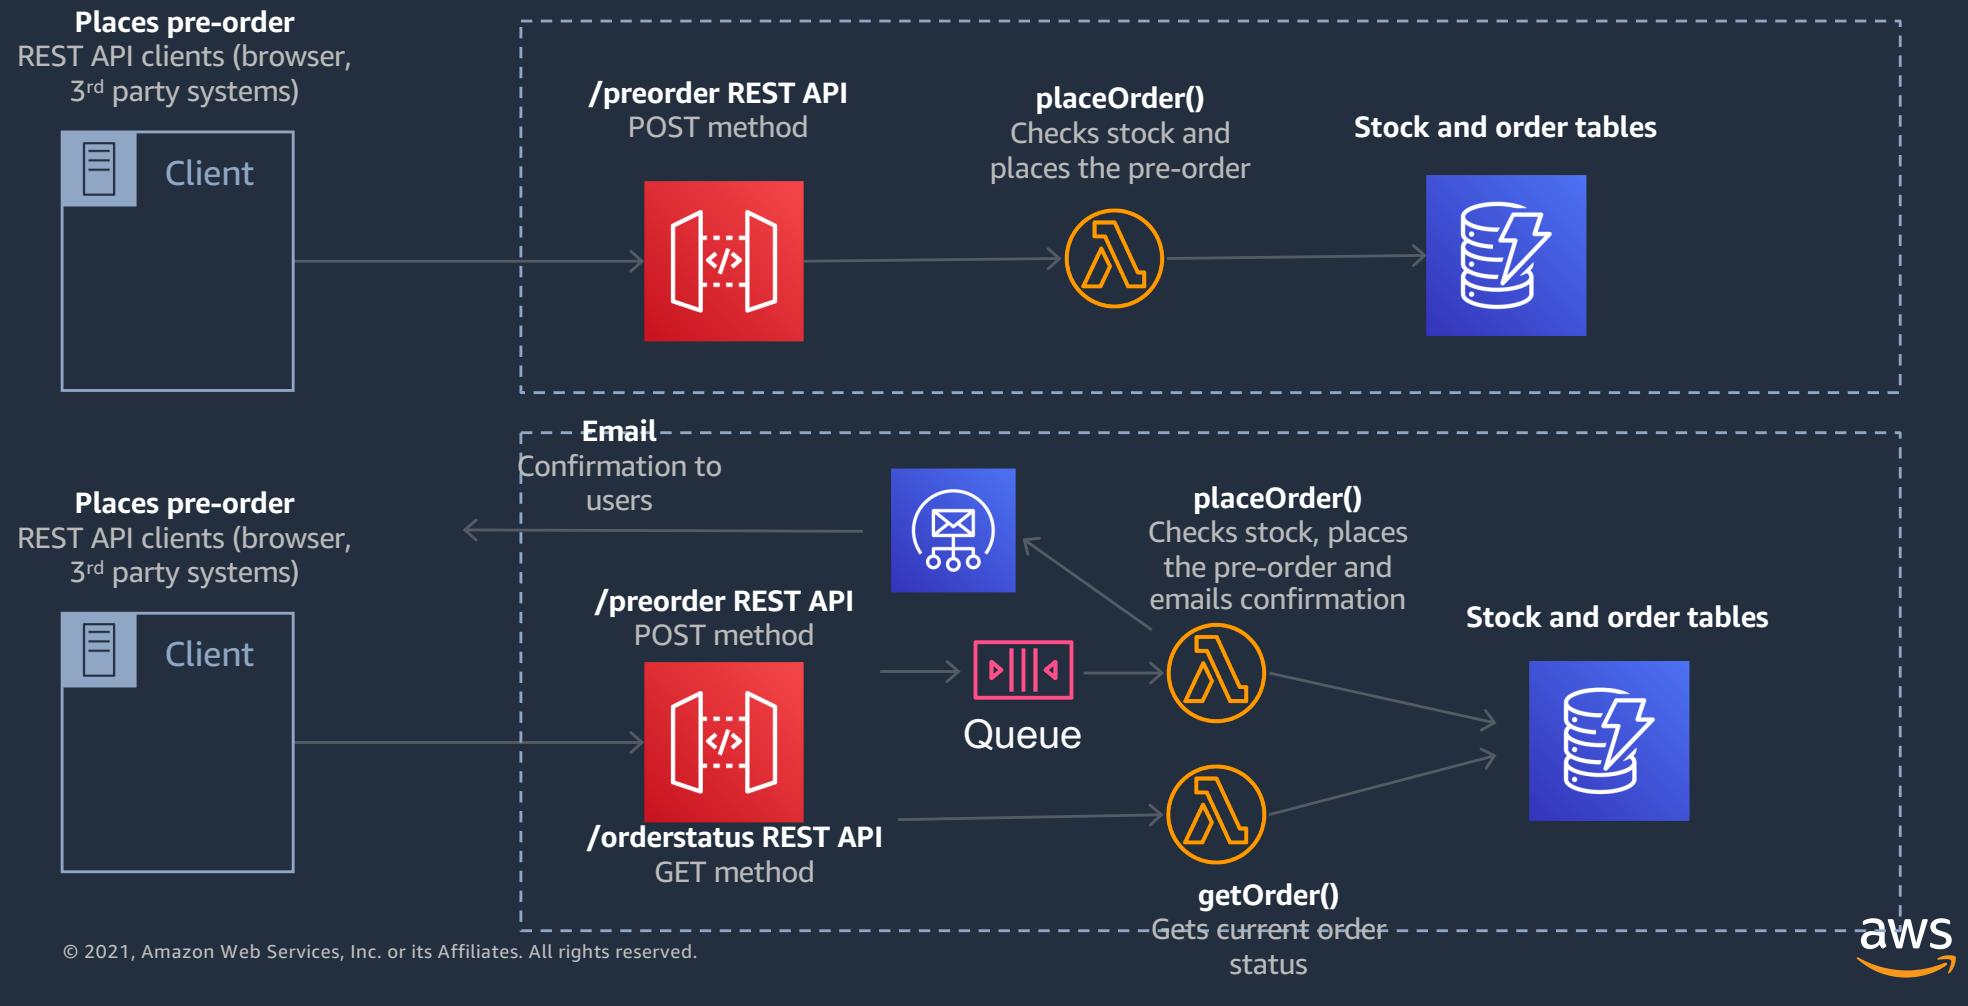


# An example serverless application

© 2021, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



# Example serverless implementation



# Is it well-architected?



## Poll#8

- Which is NOT a pillar of the AWS Well-Architected Framework?
- A: Operational Excellence
- B: User Experience
- C: Security
- D: Cost Optimization
- B

## Poll#9

- Which service is a push-based, "pub-sub" messaging service, used to "fan out" applications?
  - A: Amazon Simple Queue Service (SQS)
  - B: Amazon API Gateway
  - C: Amazon Simple Notification Service (SNS)
  - D: AWS Step Functions

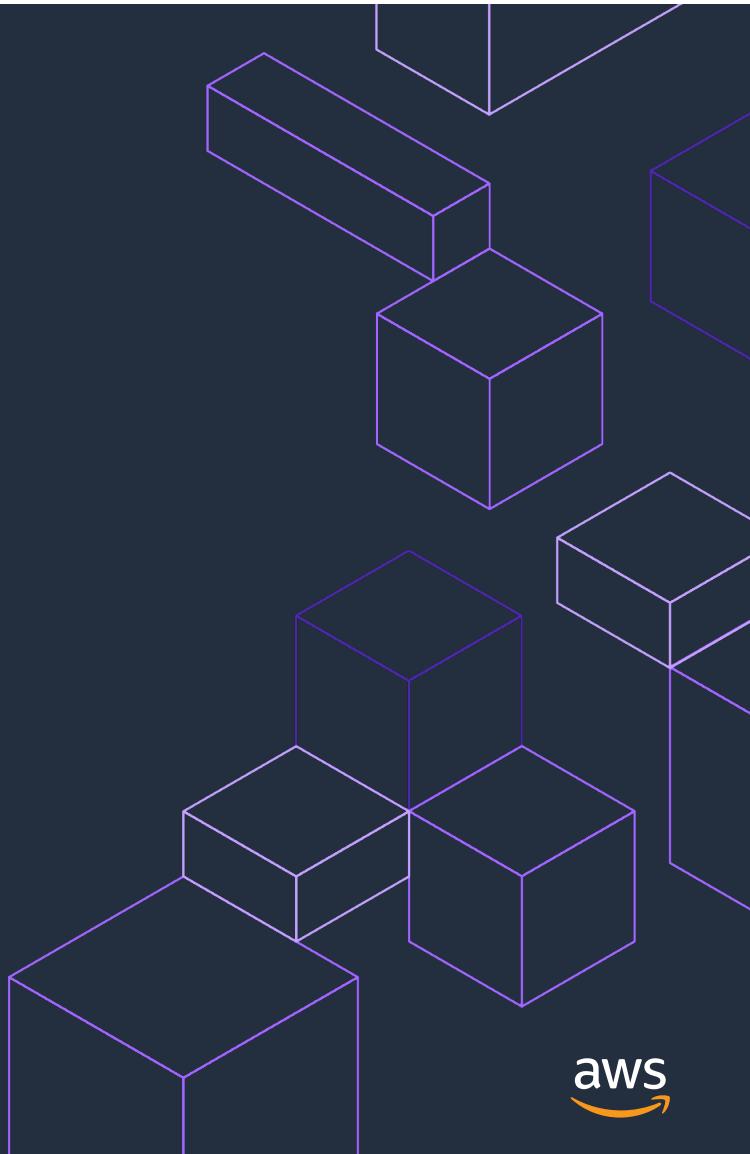
## Poll#10

- Your customer is building an internal application that serves as a repository for images uploaded by a couple of users. Whenever a user uploads an image, it **would be sent to Kinesis** for processing before it is stored in an S3 bucket. Afterwards, if the upload was successful, the application will return a prompt telling the user that the upload is successful. The entire processing typically takes about **5 minutes** to finish.

Which of the following options will allow you to **asynchronously** process the request to the application from upload request to Kinesis, S3, and return reply, in the **most cost-effective** manner? (**SELECT ONE**)

- A) Use a combination of SQS to queue the requests and then asynchronously process them using On-Demand EC2 instances
- B) Use a combination of SNS to buffer the requests and then asynchronously process them using On-Demand EC2 instance
- C) Create a Lambda function that will asynchronously process the request
- D) Use a combination of Lambda and Step Functions to orchestrate service components and **aws**  asynchronously process the requests

# Q&A



## Further reading

- [AWS Lambda FAQs](#)
- [AWS Step Functions FAQs](#)
- [Amazon SQS FAQs](#)
- [Amazon SNS FAQs](#)
- [Amazon API Gateway FAQs](#)
- [Amazon SES FAQs](#)
- [Amazon ECS FAQs](#)
- [AWS Well-Architected Framework Whitepaper](#)





# Thank you!

© 2021, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

